# Get Started with the LangChain Integration



#### NOTE

This tutorial uses LangChain's Python library <sup>™</sup>. For a tutorial that uses the JavaScript library, see Get Started with the LangChain JS/TS Integration.

You can integrate Atlas Vector Search with LangChain to build LLM applications and implement retrieval-augmented generation (RAG). This tutorial demonstrates how to start using Atlas Vector Search with LangChain to perform semantic search on your data and build a RAG implementation. Specifically, you perform the following actions:

1. Set up the environment.

★ Rate this page

**★** Ask MongoDB AI

- 2. Store custom data on Atlas.
- 3. Create an Atlas Vector Search index on your data.
- 4. Run the following vector search queries:
  - Semantic search.
  - Semantic search with score.
  - Semantic search with metadata pre-filtering.
- 5. Implement <u>RAG</u> by using Atlas Vector Search to answer questions on your data.

# **Background**

LangChain is an open-source framework that simplifies the creation of <u>LLM</u> applications through the use of "chains." Chains are LangChain-specific components that can be combined for a variety of AI use cases, including <u>RAG</u>.

By integrating Atlas Vector Search with LangChain, you can use Atlas as a vector database and use Atlas Vector Search to implement <u>RAG</u> by retrieving semantically similar documents from your data. To learn more about <u>RAG</u>, see Key Concepts.

# **Prerequisites**

To complete this tutorial, you must have the following:

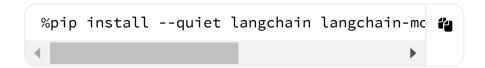
- An Atlas cluster running MongoDB version 6.0.11, 7.0.2, or later (including <u>RCs</u>). Ensure that your <u>IP address</u> is included in your Atlas project's access list.
- An OpenAl API Key. You must have a paid OpenAl account with credits available for API requests.
- An environment to run Python interactive notebooks such as Colab.<sup>™</sup>

# **Set Up the Environment**

You must first set up the environment for this tutorial. Create an interactive Python notebook by saving a file with the .ipynb extension, and then run the following code snippets in the notebook.

1 Install and import dependencies.

Run the following command:



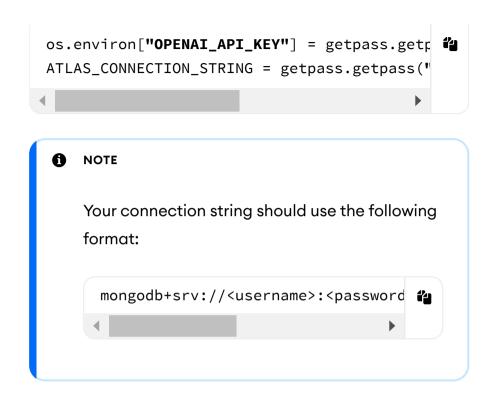
Then, run the following code to import the required packages:

import getpass, os, pymongo, pprint
from langchain\_community.document\_loaders i
from langchain\_core.output\_parsers import S
from langchain\_core.runnables import Runnak
from langchain\_mongodb import MongoDBAtlasV
from langchain\_openai import ChatOpenAI, Op
from langchain.prompts import PromptTemplat
from langchain.text\_splitter import Recursi
from pymongo import MongoClient

#### 2 Define environmental variables.

Run the following code and provide the following when prompted:

- Your OpenAl API Key.
- Your Atlas cluster's SRV connection string.



# Use Atlas as a Vector Store

Then, load custom data into Atlas and instantiate Atlas as a vector database, also called a vector store . Copy and paste the following code snippets into your notebook.

## 1 Connect to your Atlas cluster.

Run the following code to establish a connection to your Atlas cluster. It specifies the following:

- langchain\_db.test as the name of the collection for which to load the data.
- vector\_index as the name of the Atlas Vector Search index to use for querying the data.

```
# Connect to your Atlas cluster
client = MongoClient(ATLAS_CONNECTION_STRIN

# Define collection and index name
db_name = "langchain_db"
collection_name = "test"
atlas_collection = client[db_name][collectivector_search_index = "vector_index"
```

## 2 Load the sample data.

For this tutorial, you use a publicly accessible PDF document titled MongoDB Atlas Best Practices as the data source for your vector store. This document describes various recommendations and core concepts for managing your Atlas deployments.

To load the sample data, run the following code snippet. It does the following:

- Retrieves the PDF from the specified URL and loads the raw text data.
- Uses a text splitter<sup>™</sup> to split the data into smaller documents.
- Specifies chunk parameters, which determines the number of characters in each document and the number of characters that should overlap between two consecutive documents.

Document(page\_content='Mong oDB Atlas Best P ra

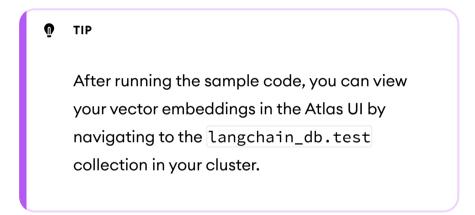
## 3 Instantiate the vector store.

Run the following code to create a vector store named vector\_store from the sample documents. This snippet uses the

MongoDBAtlasVectorSearch.from\_documents method and specifies the following parameters:

- The sample documents to store in the vector database.
- OpenAl's embedding model as the model used to convert text into vector embeddings for the embedding field.
- langchain\_db.test as the Atlas collection to store the documents.
- vector\_index as the index to use for querying the vector store.

```
# Create the vector store
vector_store = MongoDBAtlasVectorSearch.frc
    documents = docs,
    embedding = OpenAIEmbeddings(disallowed collection = atlas_collection,
    index_name = vector_search_index
)
```



# Create the Atlas Vector Search Index

To enable vector search queries on your vector store, create an Atlas Vector Search index on the langchain\_db.test collection.

## **Required Access**

To create an Atlas Vector Search index, you must have
Project Data Access Admin or higher access to the Atlas
project.

#### **Procedure**

- 1 In Atlas, go to the *Clusters* page for your project.
  - a. If it is not already displayed, select the organization that contains your desired project from the 

    Organizations menu in the navigation bar.
  - b. If it is not already displayed, select your desired project from the **Projects** menu in the navigation bar.
  - c. If the **Clusters** page is not already displayed, click **Database** in the sidebar.
- 2 Go to the Atlas Search page for your cluster.

You can go the Atlas Search page from the sidebar, the **Data Explorer**, or your cluster details page.

**Sidebar** Data Explorer Cluster Details

- a. In the sidebar, click **Atlas Search** under the **Services** heading.
- b. From the **Select data source** dropdown, select your cluster and click **Go to Atlas Search**.
- 3 Define the Atlas Vector Search index.
  - a. Click Create Search Index.
  - b. Under **Atlas Vector Search**, select **JSON Editor** and then click **Next**.
  - c. In the **Database and Collection** section, find the langchain\_db database, and select the test collection.
  - d. In the **Index Name** field, enter vector\_index.
  - e. Replace the default definition with the following index definition and then click **Next**.

This index definition specifies indexing the following fields in an index of the vectorSearch type:

embedding field as the vector type. The
 embedding field contains the embeddings
 created using OpenAl's

text-embedding-ada-002 embedding model.
The index definition specifies 1536 vector
dimensions and measures similarity using
cosine.

 page field as the filter type for pre-filtering data by the page number in the PDF.

4 Review the index definition and then click *Create Search Index*.

A modal window displays to let you know that your index is building.

5 Click Close to close the You're All Set! modal window and wait for the index to finish building.

The index should take about one minute to build. While it builds, the **Status** column reads **Initial Sync**. When it finishes building, the **Status** column reads **Active**.

## **Run Vector Search Queries**

Once Atlas builds your index, return to your notebook and run vector search queries on your data. The following examples demonstrate various queries that you can run on your vectorized data.

Semantic Search Semantic Search with Score Semant

The following query uses the similarity\_search method to perform a basic semantic search for the string

MongoDB Atlas security. It returns a list of documents ranked by relevance.



MongoDB.

Access Data Lake

- Query Federated Data
- Atlas Search
- ▼ Atlas Vector Search
  - Quick Start
  - **Create Embeddings**
  - Create and Manage Indexes
  - Create and Run Queries
  - Review Deployment

[Document(page\_content='To ensure a secure system ri Document(page\_content='MongoD B Atlas team are also Document(page\_content='MongoD B.\nMongoD B Atlas ir Document(page\_content='Atlas provides encryption of

On this page

**1** TIP

#### See also:

For a full list of semantic search methods, refer to the API reference.

Background

**Prerequisites** 

Set Up the Environment

Use Atlas as a Vector Store

Create the Atlas Vector Search Index

Run Vector Search Queries

Answer Questions on Your Data Options

- ▶ Tutorials
- ▼ Al Integrations

LangChain

LangChain JS/TS

LlamaIndex

## **Answer Questions on Your Data**

This section demonstrates how to implement <u>RAG</u> in your application with Atlas Vector Search and LangChain. Now that you've used Atlas Vector Search to retrieve semantically similar documents, run the following code examples to prompt the <u>LLM</u> to answer questions based on those documents.

**Basic RAG** RAG with Filtering

This example does the following:

- Instantiates Atlas Vector Search as a retriever to query for similar documents, including the optional k parameter to search for only the 10 most relevant documents.
- Defines a LangChain prompt template to instruct the
   LLM to use these documents as context for your query.
   LangChain passes these documents to the {context}
   input variable and your query to the {question} variable.
- Constructs a chain that specifies the following:

- Atlas Vector Search as the retriever to search for documents that are used as context by the <u>LLM</u>.
- The prompt template that you constructed.
- OpenAl's chat model as the <u>LLM</u> used to generate a context-aware response.
- Prompts the chain with a sample query about Atlas security recommendations.
- Returns the <u>LLM</u>'s response and the documents used as context. The generated response might vary.

```
# Instantiate Atlas Vector Search as a retrieve
retriever = vector_store.as_retriever(
```

# Define a prompt template template = """

search\_type = "similarity", search\_kwargs = { "k": 10 }

Use the following pieces of context to answer t If you don't know the answer, just say that you {context}

Question: {question}

111111

custom\_rag\_prompt = PromptTemplate.from\_templat

llm = ChatOpenAI()

def format\_docs(docs):

return "\n\n".join(doc.page\_content for doc

# Construct a chain to answer questions on your rag\_chain = (

{ "context": retriever | format\_docs, "quest

custom\_rag\_prompt

| llm

| StrOutputParser()

```
# Prompt the chain
question = "How can I secure my MongoDB Atlas c
answer = rag_chain.invoke(question)

print("Question: " + question)
print("Answer: " + answer)

# Return source documents
documents = retriever.get_relevant_documents(qu
print("\nSource documents:")
pprint.pprint(documents)
```

★ HIDE OUTPUT

Question: How can I secure my MongoDB Atlas cluster? Answer: To secure your MongoDB Atlas cluster, you call authentication and IP address whitelisting, review to in the MongoDB Atlas dashboard, encrypt data at rest volumes, optionally configure an additional layer of data, set up global clusters on Amazon Web Services, and Google Cloud Platform, and ensure operational call appropriate instance size, storage size, and storage Additionally, consider setting up a larger number of increased protection against database downtime.

#### Source documents:

[Document(page\_content='To ensure a secure system ri Document(page\_content='MongoD B Atlas team are also Document(page\_content='All the user needs to do in c Document(page\_content='MongoD B.\nMongoD B Atlas inc Document(page\_content='You can set up global cluster Document(page\_content='Table of Contents\n1 Introduc Document(page\_content='Atlas provides encryption of Document(page\_content='Disaster Recovery\nCreated by Document(page\_content='Security\nAs with all softwar Document(page\_content='A larger number of replica nc

# **Next Steps**

To learn about additional <u>RAG</u> use-cases with Atlas Vector Search, see the following templates provided by LangChain to help you build applications:

- Advanced RAG: Parent-Document Retrieval

MongoDB also provides the following developer resources:

- Introduction to LangChain and MongoDB Atlas Vector Search
- RAG with Atlas Vector Search, LangChain, and OpenAl
- Leveraging MongoDB Atlas Vector Search with LangChain



TIP

#### See also:

- LangChain Documentation
- LangChain API Reference<sup>™</sup>

## © English

### **About**

Careers Investor Relations

Legal GitHub

Security Information Trust Center

Connect with Us

# Support

Contact Us Customer Portal

Atlas Status Customer Support

Manage Cookies

# **Deployment Options**

MongoDB Atlas

Enterprise Advanced

**Community Edition** 

© 2024 MongoDB, Inc.